


G. Mikołaj Boć 

# WASM - doing C++ the web way



# Agenda

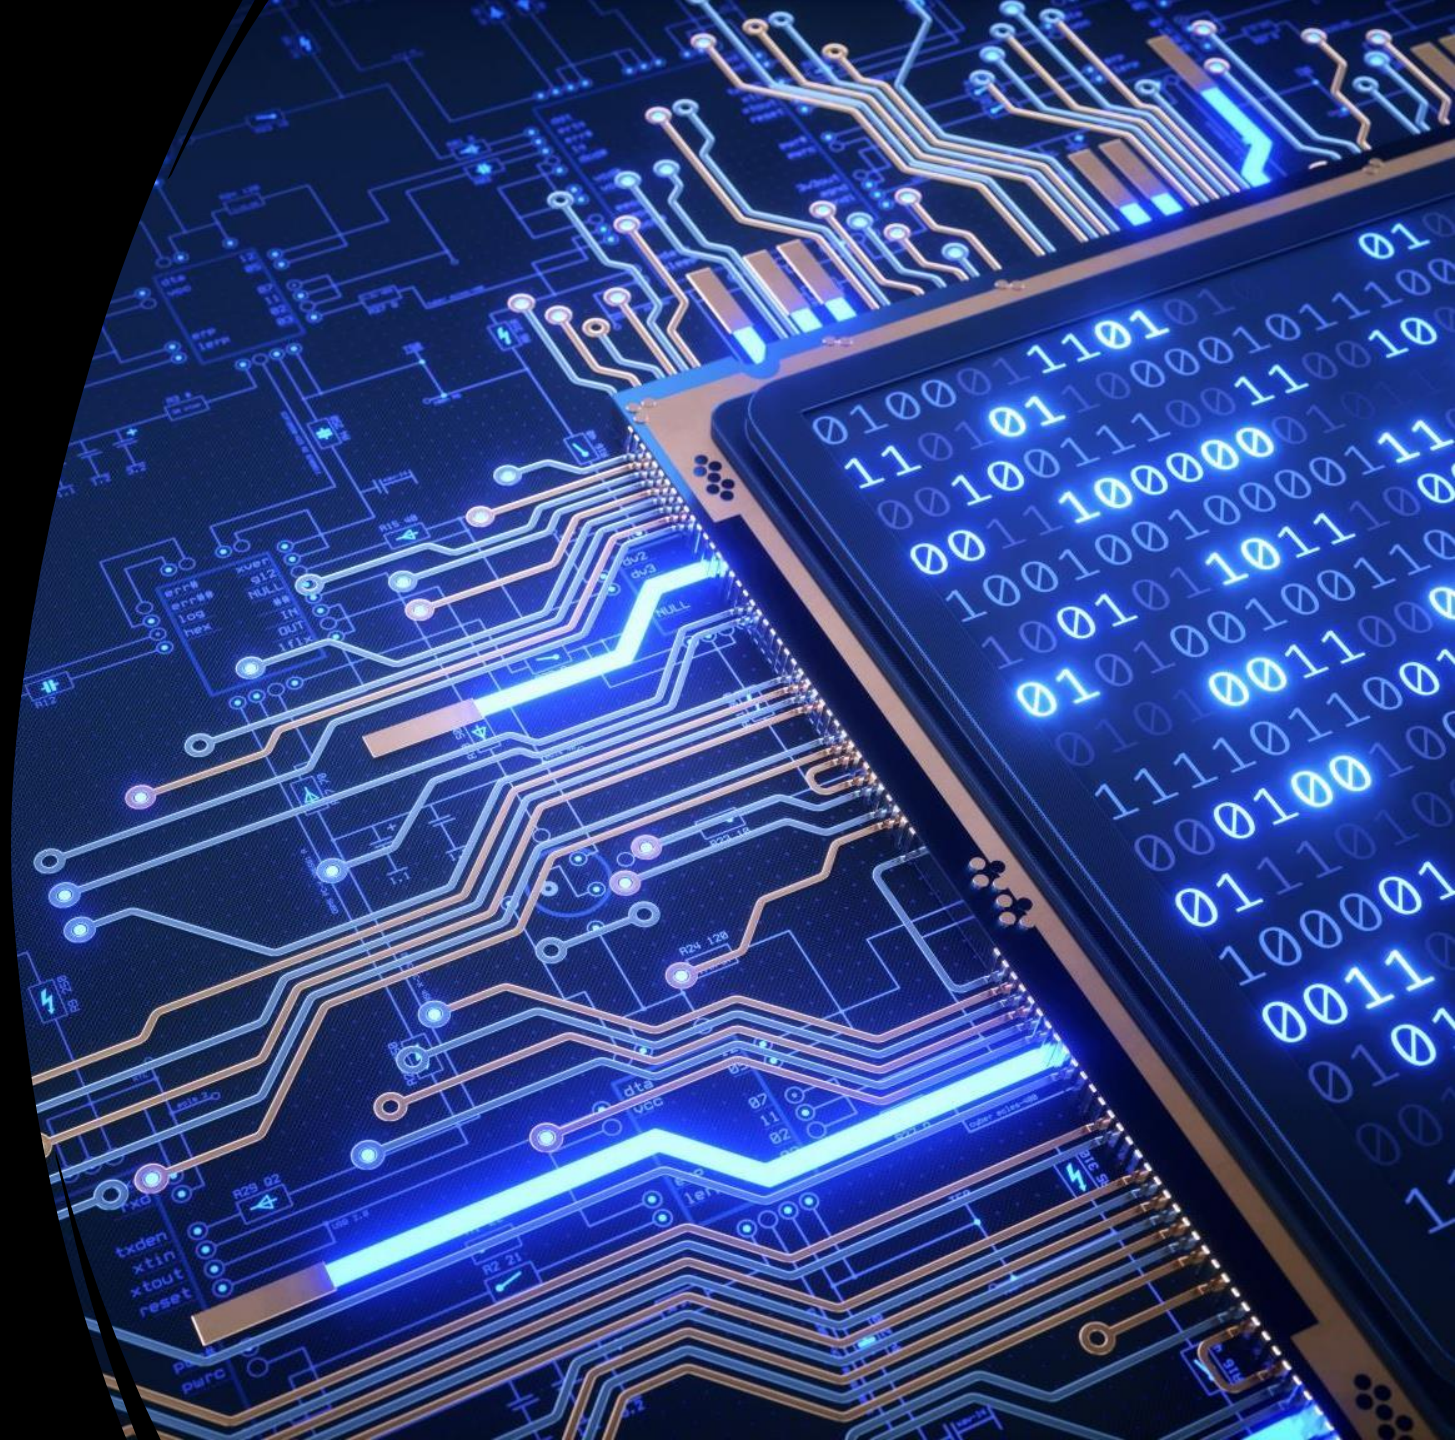
- What is it?
- Where/how can I run it?
- How fast is it?
- What are the potential challenges?



# What is it?

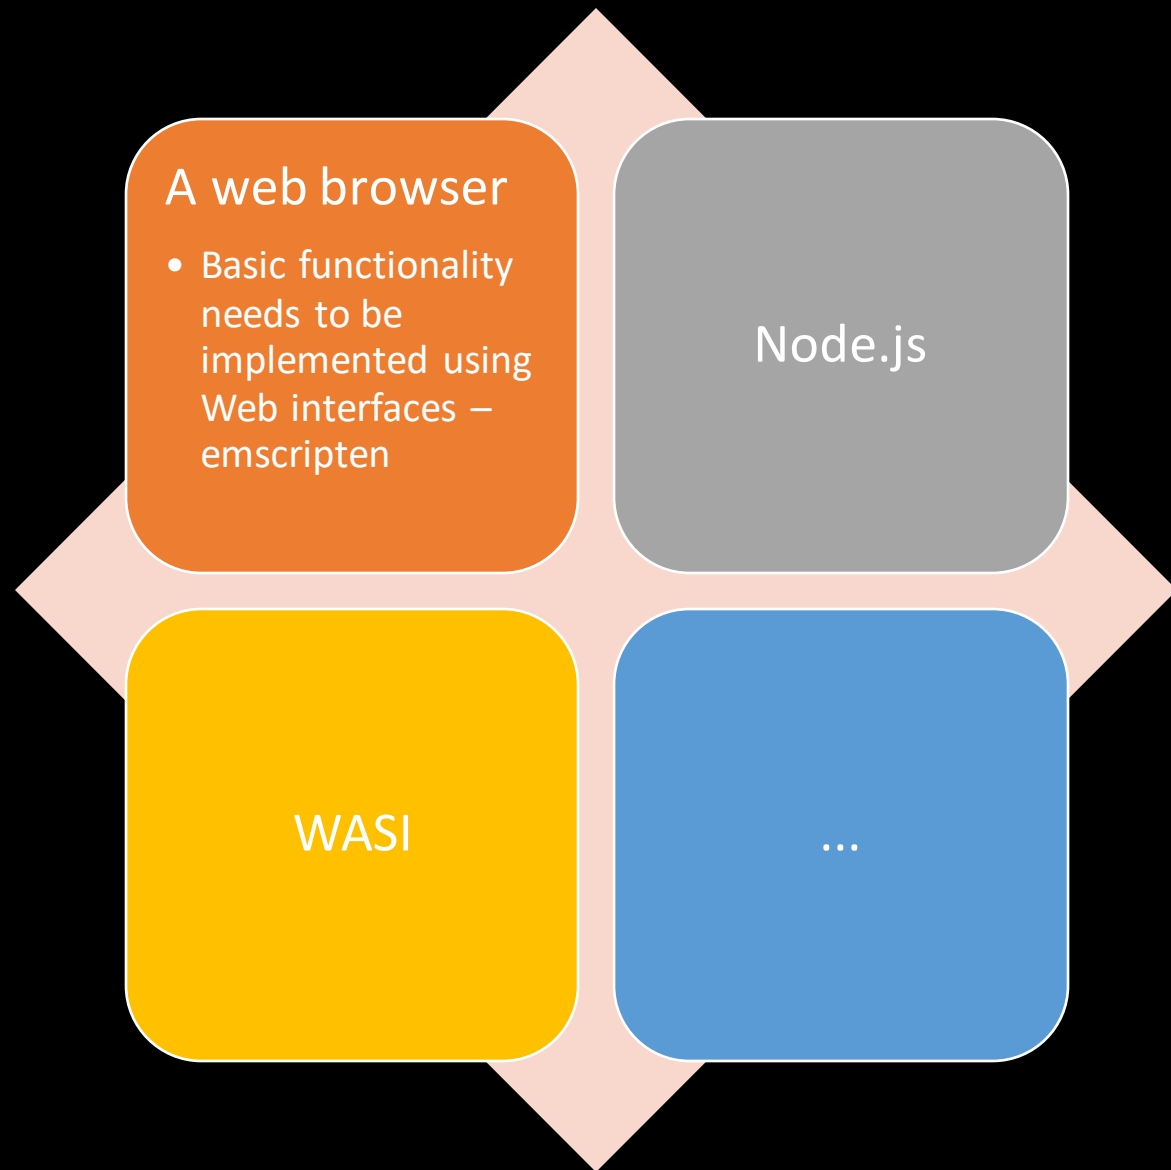
---

- Binary instruction format
- Portable – virtual machine
- Some call it the holy grail of portability
- Tiered compilation to machine code, e.g., Chromium's **Liftoff** + **TurboFan**
- Multitude of languages supported via various toolchains (C++ among them)
- With little extra effort, as debuggable as JS on the web



# Where can I run it?

*"Its main goal is to enable high performance applications on the Web, but it does not make any Web-specific assumptions or provide Web-specific features, so it can be employed in other environments as well"*





WasmBoy (Web Assembly, Assemblyscript)

Frames Run: 2500

Current FPS Average: 554



WasmBoy (Typescript)

Frames Run: 2500

Current FPS Average: 475



WasmBoy (Typescript, Closure Compiled)

Frames Run: 2500

Current FPS Average: 526



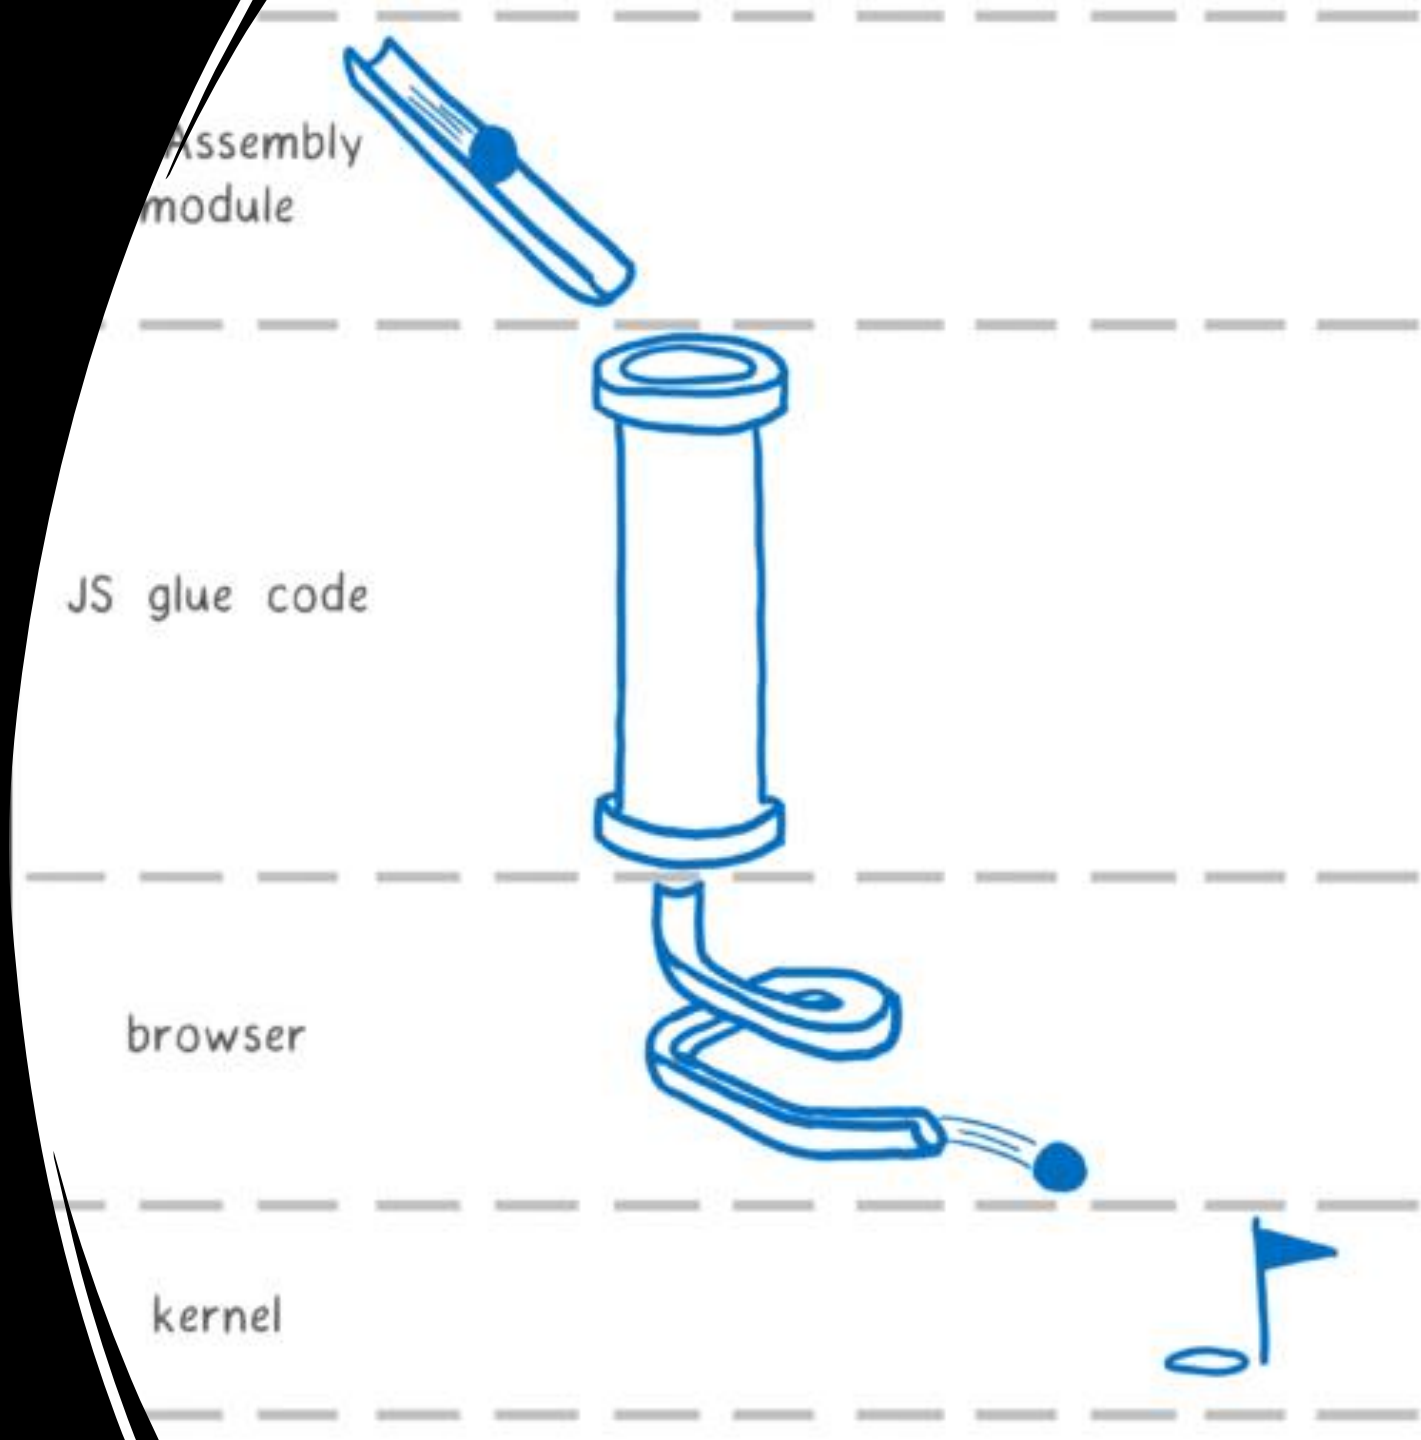
# How fast is it?

- Vs JavaScript, almost always positive performance impact, sometimes substantial (7% - 50% faster in various studies)
- Vs native, various studies
  - Small, focused scientific programs ~10% slower
  - Larger programs up to 50%-250% slower

# emscripten

- Complete toolchain to compile C/C++ to WASM (LLVM + Binaryen + Closure and others)
- Implementation of system APIs
  - File access
  - pthreads using web workers
  - Emulated POSIX TCP sockets
  - OpenAL with Web Audio
  - Wrappers for use in C++ codebase
    - `emscripten_fetch`, `emscripten_set_X_callback`
  - ...
- Inline JS support
- JS-to-C++ transliteration via `emscripten::val`
- Convenience code for using web APIs
  - WebGL

# emscripten



# What are the potential challenges?

---

- Permissions – sandbox in renderer process
- File access
- The async vs sync dilemma
- Linking – system libraries
- Aggressive symbol removal
- Executing apps for output
- Starving the main loop
- Threading <-> web workers
- Copying data between memory areas



# Permissions



## **Accessing certain APIs requires special permissions**

Video recording, audio recording, notifications, clipboard access etc.



## **Need to think when to request the permissions**

Transient activation needed for requesting permissions



# File access

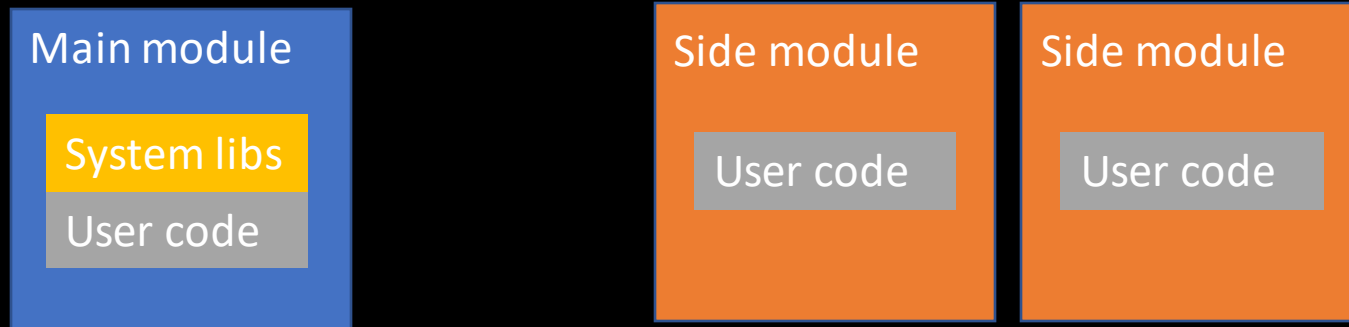
- File access is heavily limited
- Various modes that emulate a filesystem
  - MEMFS
  - NODEFS
  - IDBFS – has to be synchronized so a bit of an overhead
  - ...
- Can also access the filesystem via the filesystem API
  - `show(?:Open|Save)FilePicker` – only available on Chrome and needs OPFS/permissions and transient activation

# Async operations on the web

- Mostly async operations on the web – promises and async callbacks to events
  - Sync emulation in emscripten - hogging the main event loop
  - `emscripten_set_main_loop`
- Starting an event loop is a challenge
  - Asyncify
  - Moving the main thread to a worker
- `emscripten::val` + callbacks
  - Promise is particularly problematic

# Linking

- Apps need to ship system library implementations
- Dynamic:
  - Side modules + main modules



- Load time
  - Caching helps with common libraries

# Aggressive symbol removal

- Tries to keep the binary small
- Dynamically calling system functions that are deemed unused in the binary is impossible if the function is not explicitly marked for inclusion

```
29 // Volatile is to make this unoptimizable, so that the function is referenced, but is not
30 // called at runtime.
31 volatile bool doIt = false;
32 if (doIt)
33     emscripten_set_wheel_callback(NULL, 0, 0, NULL);
```



# Executing apps

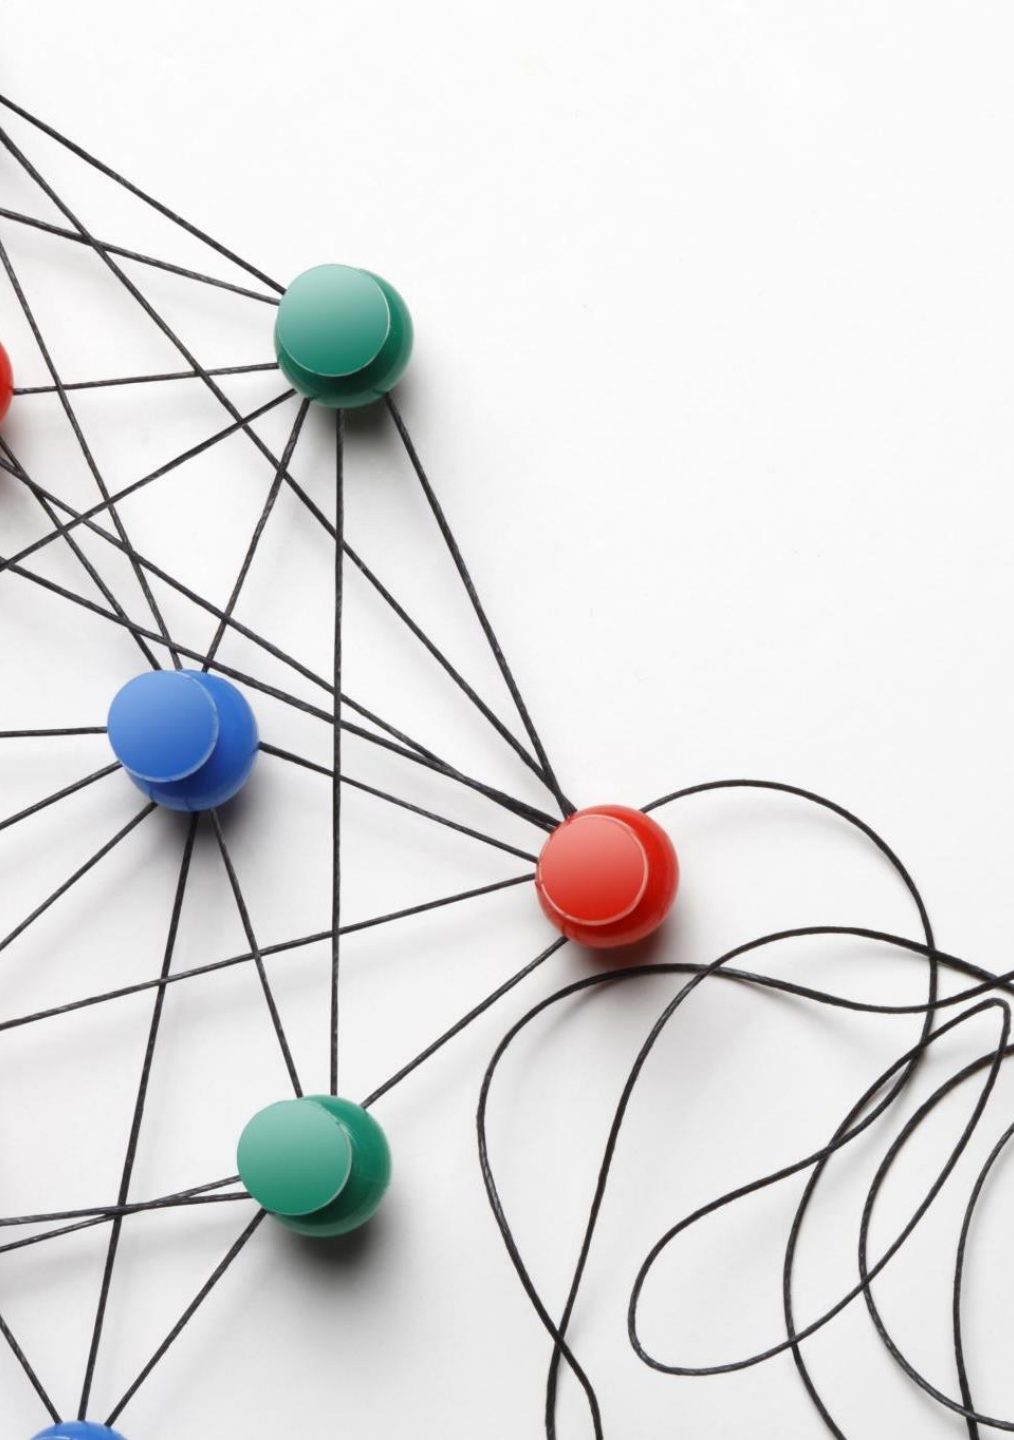
- Simply loading the generated html file
  - Cannot make the program a part of an execution pipeline
  - Difficult to get the output
- Using web APIs + JS to load the module
  - Most control
  - More work
  - emscripten quirks – more complicated functionality needs a lot of code reading
- emrun
  - Sets up a local server which reads POST data sent from the application
  - Application's output redirected to the server using a js preamble file
  - Problems:
    - Selecting a browser
    - Port clashes

# Starving the main loop

---

- When running long operations, it is necessary to yield to the main event loop
- There's an option to move the main application thread to a secondary thread to avoid such problems
  - Target of experiments in Qt now





# Threading

- Thread implementation needs SharedArrayBuffer to work
- Extra worker file generated
- Quirks
  - Proxying on child threads
    - Can lead to deadlocks!
  - Busy-waiting for `pthread_join`, `pthread_mutex_lock`, `usleep`, etc.
  - `pthread_create` needs to return to the main event loop
    - `PTHREAD_POOL_SIZE` – if you expect to run a certain number of threads at a time
    - `PROXY_TO_PTHREAD` – runs main on a worker. Introduces other problems

# Memory areas

- Emscripten heap
- JS memory
- Might need to copy memory around to get the benefits of WASM
  - Data from direct calls to javascript
- Copying not always necessary –  
`typed_memory_view`

Thank you!



# It doesn't end

- By default, runtime does not exit
  - Stdio streams not flushed
  - Global destructors not called
  - atexit callbacks not called
- -s EXIT\_RUNTIME